CMSC 25025 / STAT 37601

Machine Learning and Large Scale Data Analysis

Assignment 2

Due: Tuesday, April 17, 2017 at 1:30 pm.

Please hand in this Homework in 5 files:

1. A pdf of the theoretical homework (problems 1-2).
2. A pdf of your jupyter notebook for problem 3.
3. The ipynb file for problem 3.
4. A pdf of your jupyter notebook for problem 4.
5. The ipynb file for problem 4.

1. *PCA* (10 points)

   The problem of fitting the best k-dimensional subspace to data $x_1, \ldots, x_n \in \mathbb{R}^d$ can be written as the optimization

   $$\min_{\mu, \{\lambda_i\}, V_k} \sum_{i=1}^{n} \|x_i - \mu - V_k \lambda_i\|^2$$

   where V_k is an $d \times k$ orthogonal matrix. Show that an optimum over the variables $\mu \in \mathbb{R}^d$ and $\lambda_i \in \mathbb{R}^k$ is given by

   $$\widehat{\mu} = \overline{x}_n = \frac{1}{n} \sum_{i=1}^{n} x_i$$

   $$\widehat{\lambda}_i = V_k^T (x_i - \overline{x}_n).$$

   Show that $\widehat{\mu}$ is not unique, and characterize the set of possible solutions.

2. *Lasso* (10 points)

   Show that the minimizer of the function $f(\beta) = -t\beta + \beta^2 + \lambda|\beta|$ is given by $sign(t)[t - \lambda]_+$

3. *Analysis of SOU similarities* (20 points)

   In this problem you will use the classic *vector space model* from information retrieval to find similar SOU addresses. Your code from Assignment 1 may come in handy here.

   In the vector space model, a document of words d is represented by a TF-IDF vector $\mathbf{w}(d) = (w_1(d), w_2(d), \ldots, w_V(d))$ of length V, where V is the total number of words

in the vocabulary. The TF-IDF weights are given by

$$w_i(d) = n_i(d) \cdot \log\left(\frac{|\mathcal{D}|}{\sum_{d' \in \mathcal{D}} \mathbb{1}(t_i \in d')} \right)$$

where $n_i(d)$ is the number of times term t_i appears in document d, $\sum_{d' \in \mathcal{D}} \mathbb{1}(t_i \in d')$ is the number of documents that contain term t_i, and $|\mathcal{D}| = \sum_{d' \in \mathcal{D}} 1$ is the total number of documents in the collection $\mathcal{D}$. This weighting scheme favors terms that appear in few documents.

(a) Compute the TF-IDF vectors for each SOU address. You should lower case all of the text, and remove punctuation. For example, you could use something like this:

```
def clean_and_split(s):
        # encode to UTF-8, convert to lowercase and translate all hyphens and
        # punctuation to whitespace
        s = s.encode('utf-8').lower().replace('-',' ').translate(None, string.punctuation)
        # replace \r\n
        s = re.sub('(\r\n)+',' ',s)
        # replace whitespace substrings with one whitespace and remove
        # leading/trailing whitespaces
        s = re.sub(' +',' ',s.strip())
        return s.split(' ')
```

You will have to make choices about the size of the term vocabulary to use—for example throwing out the 20 most common words, and words that appear fewer than, say, 50 times.

(b) A similarity measure between documents is

$$\text{sim}(d, d') = \frac{\mathbf{w}(d) \cdot \mathbf{w}(d')}{\|\mathbf{w}(d)\| \|\mathbf{w}(d')\|},$$

the cosine of the angle between the corresponding TF-IDF vectors. In terms of this measure, find the

- 50 most similar pairs of SOUs given by different Presidents.
- 50 most similar pairs of SOUs given by the same President.
- 25 most similar pairs of *Presidents*, averaging the cosine similarity over all pairs of their SOUs.

When you read the above speeches, do they indeed seem similar to you? (You can read the speeches in a more reader-friendly format here: `http://www.presidency.ucsb.edu/sou.php`) Comment on what you find, and describe what is needed to construct a better similarity measure between documents.

(c) Using this vector representation, cluster the speeches using k-means. To do this you can use the `sklearn` implementation of k-means:

```
from sklearn.cluster import KMeans
```

then we can run the method `KMeans.train`:

```
model = KMeans(n_clusters=c)
sou_clust=model.fit(tf_idf_mat)
```

The options here limit the number of iterations of kmeans to 50, the number of clusters to 10, the clusters are initialized randomly.

Experiment with different number of clusters, and display the clusters obtained (in some manner that you choose). Comment on the clustering results, and whether or not the results are interpretable.

4. *Handwritten digits* (60 points)

For this extended problem you are going to work with the MNIST handwritten digits dataset. The dataset is made up of more than 70,000 images of the digits 0-9. The data are length 784 vectors ranging between [0,255] representing intensity values. The vectors are a flattened version of the 28×28 pixel scans of the digits. Each data point also has a response value, corresponding to which digit it is. Here is an example of how to read the data.

```
import numpy as np
data=np.float64(np.load('/project/cmsc25025/mnist/MNIST.npy'))
labels=np.float32(np.load('/project/cmsc25025/mnist/MNIST_labels.npy'))
```

After you have read in the data, you should divide the data by 255 so that the values lie in [0, 1]. Divide the data into 60% training, 20% development, 20% test.

To display the images you can do something like this: store a few digits into a numpy array `subset`, whose rows are digits.

```
%matplotlib inline
import matplotlib.pyplot as plt

nrows = 4
ncols = 5
plt.figure(figsize=(ncols*2, nrows*2))

for i in xrange(nrows*ncols):
    plt.subplot(nrows, ncols, i+1)
    plt.imshow(subset[i].reshape((28,28)), cmap='gray')
    plt.axis('off')

plt.axis('off')
plt.show()
```

This problem has four subparts: (1) PCA, (2) k-means, (3) spectral clustering, and (4) classification.

Part 1: PCA

(a) *Extract principal components.*

Perform PCA to extract the principal components of the training data. You may not use a library to do PCA for you, but you may use libraries (for example from numpy) to compute the singular value or eigenvalue decomposition of a matrix. Display the first 10 principal components as images.

(b) *Plot variance*

Plot the variance of all of the principal components—this corresponds to the singular values. This should be monotonically decreasing.

(c) *Dimension reduction*

Take a data point in the test data set and project it onto the first m principal components (remember to first subtract the mean). Then, transform that m-length vector back into a 784-length vector and display it as an image (add back the mean). Repeat this for several different values of m. Also, try it on different data points. Describe the results qualitatively. Does it give an accurate representation of the images? How do the results depend on m? Can you describe what the top principal components are capturing?

Part 2: k-means

Now use k-means to perform unsupervised clustering of the digit data, and try to assess how well the clusters capture the structure of the data. Do not produce more than 50 clusters. Show the centers of the clusters. To evaluate how closely the clusters capture the structure of the data, associate each cluster with the majority label.

Construct plots showing samples of the digits from each cluster, and comment on how well the clusters respect the true digit labels.

Part 3: Spectral clustering

Spectral clustering uses k-means on top of a low dimensional representation of data. It makes use of the spectrum (eigenvectors) of a particular similarity matrix of the data in order to perform the dimensionality reduction.

Construct the similarity matrix A by applying Gaussian kernel to squared distance:

$$W_{ij} = \exp(-\|x_i - x_j\|^2/h),$$

where x_i, x_j are digits. The value of the bandwidth h is crucial for good clustering performance. You should experiment with different values to get an appropriate scaling. Next, compute the normalized graph Laplacian matrix

$$L = I - \widetilde{W} = I - D^{-1/2}WD^{-1/2}$$

where D is the diagonal matrix $D_{ii} = \sum_j W_{ij}$.

Compute the bottom few eigenvectors of L except the lowest one. Read off the rows $v_{i1}, v_{i,2}, ..v_{ir}$ as the embedding of the digits. Run k-means. Typical choices of the reduced dimension are r=2 or r=3. Compare the result with standard k-means, and describe your findings.

Part 4: Classification

Next you will explore how well logistic regression on the raw data compares with logistic regression on the PC's for classification.

(a) *Multinomial logistic regression* is a linear model for multi-class classification. The conditional probabilities of the outcome class $k \in [K]$ are modeled using the softmax function

$$\mathbb{P}(y = k|x, \beta_k, b_k) = \frac{\exp(\beta_k^\top x + b_k)}{\sum_{j=1}^{K} \exp(\beta_j^\top x + b_j)}.$$

Using the noramlized data from part 2, train the model on the training set, evaluate the model on development set, and calculate the error rate. You will need the `sklearn.linear_model` implementation of multiclass logistic regression:

```
from sklearn.linear_model import LogisticRegression
lg=LogisticRegression(fit_intercept=True, C=100000, penalty='l2',
          multi_class='muiltinomial',solver='lbfgs')
```

This implementation of logistic regression always assumes a regularization term and $C = 1/\lambda$. So small regulatization term corresponds to very large C.

The next step is to determine how well PCA works in terms of classification. Use the PC's you computed earlier. Train a classifier to predict the class label given those principal components, using logistic regression. To predict the label for a new data point, we use its projection onto the principal eigenvectors.

After setting up the basic work flow you can use cross-validation to select the number of principal components.

(b) Train multinomial logistic regression model on training data, using the top k principal components. Predict the labels of the development data and compute the error rate. Plot the error as a function of k. How does the error change as the number k vary?

(c) Pick k that minimizes the error on the development data. Retrain the model using both training and development data. Compute the error rate on testing data. Compare the result to multinomial logistic regression using raw features, and describe your findings.